

ETC5521: Exploratory Data Analysis

Initial data analysis

Lecturer: *Emi Tanaka*

✉ ETC5521.Clayton-x@monash.edu

📅 Week 3 - Session 1



Data Analysis



Data analysis is a process of cleaning, transforming, inspecting and modelling data with the aim of extracting information.

- Data analysis includes:
 - exploratory data analysis,
 - confirmatory data analysis, and
 - *initial data analysis*.
- Confirmatory data analysis is focussed on statistical inference and includes processes such as testing hypothesis, model selection, or predictive modelling... but today's focus will be on ***initial data analysis***.

Initial Data Analysis (IDA)

- There are various definitions of IDA, much like there are numerous definitions for EDA.
- Some people would be practicing IDA without realising that it is IDA.
- Or other cases, a different name is used to describe the same process, such as Chatfield (1985) referring to IDA also as **"initial examination of data"** and Cox & Snell (1981) as **"preliminary data analysis"** and Rao (1983) as **"cross-examination of data"**.

So what is IDA?

What is IDA?



The two **main objectives for IDA** are:

1. **data description**, and
2. **model formulation**.

- ***IDA differs from the main analysis*** (i.e. usually fitting the model, conducting significance tests, making inferences or predictions).
- ***IDA is often unreported*** in the data analysis reports or scientific papers due to it being "uninteresting" or "obvious".
- The role of ***the main analysis is to answer the intended question(s) that the data were collected for***.
- Sometimes IDA alone is sufficient.

1 Data Description Part 1/2

- Data description should be one of the first steps in the data analysis to ***assess the structure and quality of the data***.
- We refer them to occasionally as ***data sniffing*** or ***data scrutinizing***.
- These include using common or domain knowledge to check if the recorded data have sensible values. E.g.
 - Are positive values, e.g. height and weight, recorded as positive values with a plausible range?
 - If the data are counts, do the recorded values contain non-integer values?
 - For compositional data, do the values add up to 100% (or 1)? If not is that a measurement error or due to rounding? Or is another variable missing?

1 Data Description Part 2/2

- In addition, numerical or graphical summaries may reveal that there is unwanted structure in the data. E.g.,
 - Does the treatment group have different demographic characteristics to the control group?
 - Does the distribution of the data imply violations of assumptions for the main analysis?
- *Data sniffing* or *data scrutinizing* is a process that you get better at with practice and have familiarity with the domain area.
- Aside from checking the *data structure* or *data quality*, it's important to check how the data are understood by the computer, i.e. checking for *data type* is also important. E.g.,
 - Was the date read in as character?
 - Was a factor read in as numeric?

Next we'll see some *illustrative* *examples* and *cases* based on real data with some R codes

- Note: that there are a variety of ways to do IDA & EDA and you don't need to prescribe to what we show you.

Example 1 Checking the data type Part 1/2

lecture3-example.xlsx

	A	B	C	D
1	id	date	loc	temp
2	1	3/1/10	New York	42
3	2	3/2/10	New York	41.4
4	3	3/3/10	New York	38.5
5	4	3/4/10	New York	41.1
6	5	3/5/10	New York	39.8

```
library(readxl)
library(here)
df <- read_excel(here("data/lecture3-example.xlsx"))
df

## # A tibble: 5 x 4
##       id date                loc      temp
##   <dbl> <dtm>                <chr>    <dbl>
## 1     1 2010-01-03 00:00:00 New York 42
## 2     2 2010-02-03 00:00:00 New York 41.4
## 3     3 2010-03-03 00:00:00 New York 38.5
## 4     4 2010-04-03 00:00:00 New York 41.1
## 5     5 2010-05-03 00:00:00 New York 39.8
```

Any issues here?

Example 1 Checking the data type Part 2/2

```
library(lubridate)
df %>%
  mutate(id = as.factor(id),
         day = day(date),
         month = month(date),
         year = year(date)) %>%
  select(-date)
```

```
## # A tibble: 5 x 6
##   id      loc      temp    day month  year
##   <fct> <chr>    <dbl> <int> <dbl> <dbl>
## 1 1      New York  42      3      1  2010
## 2 2      New York  41.4     3      2  2010
## 3 3      New York  38.5     3      3  2010
## 4 4      New York  41.1     3      4  2010
## 5 5      New York  39.8     3      5  2010
```

- `id` is now a **factor** instead of **integer**
- `day`, `month` and `year` are now extracted from the `date`
- Is it okay now?
- In the United States, it's common to use the date format MM/DD/YYYY (gasps) while the rest of the world commonly use DD/MM/YYYY or YYYY/MM/DD.
- It's highly probable that the dates are 1st-5th March and not 3rd of Jan-May.
- You can validate this with other variables, say the temperature [here](#).

Example 1 Checking the data type with R Part 1/3

- You can robustify your workflow by ensuring you have a check for the expected data type in your code.

```
xlsx_df <- read_excel(here("data/lecture3-example.xlsx"),  
                      col_types = c("text", "date", "text", "numeric")) %>%  
  mutate(id = as.factor(id),  
         date = as.character(date),  
         date = as.Date(date, format = "%Y-%d-%m"))
```

- `read_csv` has a broader support for `col_types`

```
csv_df <- read_csv(here("data/lecture3-example.csv"),  
                  col_types = cols(  
    id = col_factor(),  
    date = col_date(format = "%m/%d/%y"),  
    loc = col_character(),  
    temp = col_double()))
```

- The checks (or coercions) ensure that even if the data are updated, you can have some confidence that any data type error will be picked up before further analysis.

Example 1 Checking the data type with R Part 2/3

You can have a quick glimpse of the data type with:

```
dplyr::glimpse(xlsx_df)
```

```
## Rows: 5  
## Columns: 4  
## $ id    <fct> 1, 2, 3, 4, 5  
## $ date  <date> 2010-03-01, 2010-03-02, 2010-03-03, 2010-03-04, 2010-03-05  
## $ loc   <chr> "New York", "New York", "New York", "New York", "New York"  
## $ temp  <dbl> 42.0, 41.4, 38.5, 41.1, 39.8
```


```
dplyr::glimpse(csv_df)
```

```
## Rows: 5  
## Columns: 4  
## $ id    <fct> 1, 2, 3, 4, 5  
## $ date  <date> 2010-03-01, 2010-03-02, 2010-03-03, 2010-03-04, 2010-03-05  
## $ loc   <chr> "New York", "New York", "New York", "New York", "New York"  
## $ temp  <dbl> 42.0, 41.4, 38.5, 41.1, 39.8
```


Example 1 Checking the data type with R Part 3/3

You can also visualise the data type with:

```
library(visdat)  
vis_dat(xlsx_df)
```



```
library(inspectdf)  
inspect_types(xlsx_df) %>%  
  show_plot()
```



Example 2 Checking the data quality

```
df2 <- read_csv(here("data/lecture3-example2.csv"),
  col_types = cols(id = col_factor(),
                    date = col_date(format = "%m/%d/%y"),
                    loc = col_character(),
                    temp = col_double()))
```

df2

```
## # A tibble: 9 x 4
##   id    date      loc      temp
##   <fct> <date>    <chr>    <dbl>
## 1 1      2010-03-01 New York    42
## 2 2      2010-03-02 New York   41.4
## 3 3      2010-03-03 New York   38.5
## 4 4      2010-03-04 New York   41.1
## 5 5      2010-03-05 New York   39.8
## 6 6      2020-03-01 Melbourne  30.6
## 7 7      2020-03-02 Melbourne  17.9
## 8 8      2020-03-03 Melbourne  18.6
## 9 9      2020-03-04 <NA>      21.3
```

- Numerical or graphical summaries or even just eye-balling the data helps to uncover some data quality issues.
- Any issues here?
- There's a missing value in `loc`.
- Temperature is in Fahrenheit for New York but Celsius in Melbourne (you can validate this again using external sources).

Case study 1 Soybean study in Brazil Part 1/3

```
data("lehner.soybeanmold", package = "agridat")
skimr::skim(lehner.soybeanmold)
```


```
## — Data Summary —————
##                               Values
## Name                         lehner.soybeanmold
## Number of rows                382
## Number of columns             9
## -----
## Column type frequency:
##   factor                       4
##   numeric                      5
## -----
## Group variables               None
##
```

scroll



Case study 1 Soybean study in Brazil Part 2/3


```
vis_miss(lehner.soybeanmold)
```



```
inspect_na(lehner.soybeanmold) %>%  
show_plot()
```

Prevalence of NAs in df::lehner.soybeanmold

df::lehner.soybeanmold has 9 columns, of which 1 have missing values

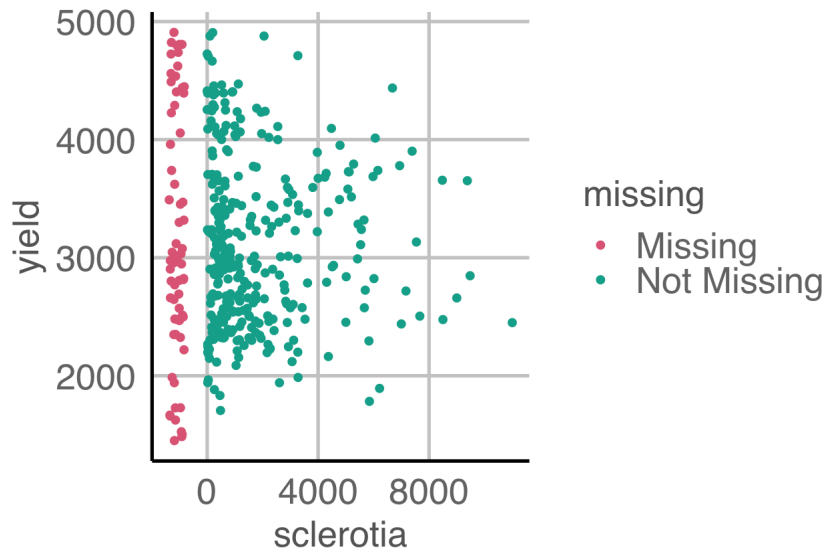


Case study 1 Soybean study in Brazil

Part 3/3

Checking if missing values have different yields:

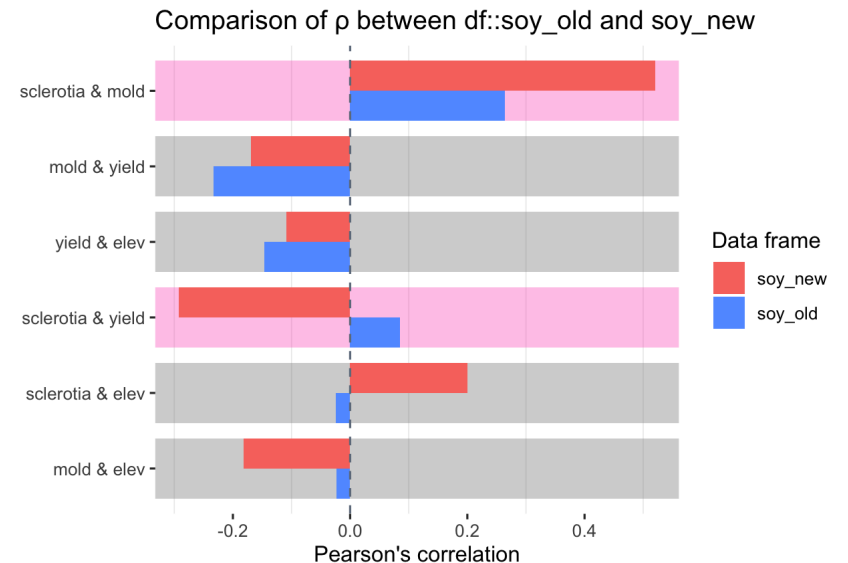
```
library(naniar)
ggplot(lehner.soybeanmold,
       aes(sclerotia, yield)) +
  geom_miss_point() +
  scale_color_discrete_qualitative()
```



Compare the new with old data:

```
soy_old <- lehner.soybeanmold %>%
  filter(year %in% 2010:2011)
soy_new <- lehner.soybeanmold %>%
  filter(year == 2012)
```

```
inspect_cor(soy_old, soy_new) %>%
  show_plot()
```



Sanity check your data

Case study 2 Employment Data in Australia Part 1/3

Below is the data from ABS that shows the total number of people employed in a given month from February 1976 to December 2019 using the original time series.

```
glimpse(employed)
```

```
## Rows: 509
```

```
## Columns: 4
```

```
## $ date <date> 1978-02-01, 1978-03-01, 1978-04-01, 1978-05-01, 1978-06-01, 19
```

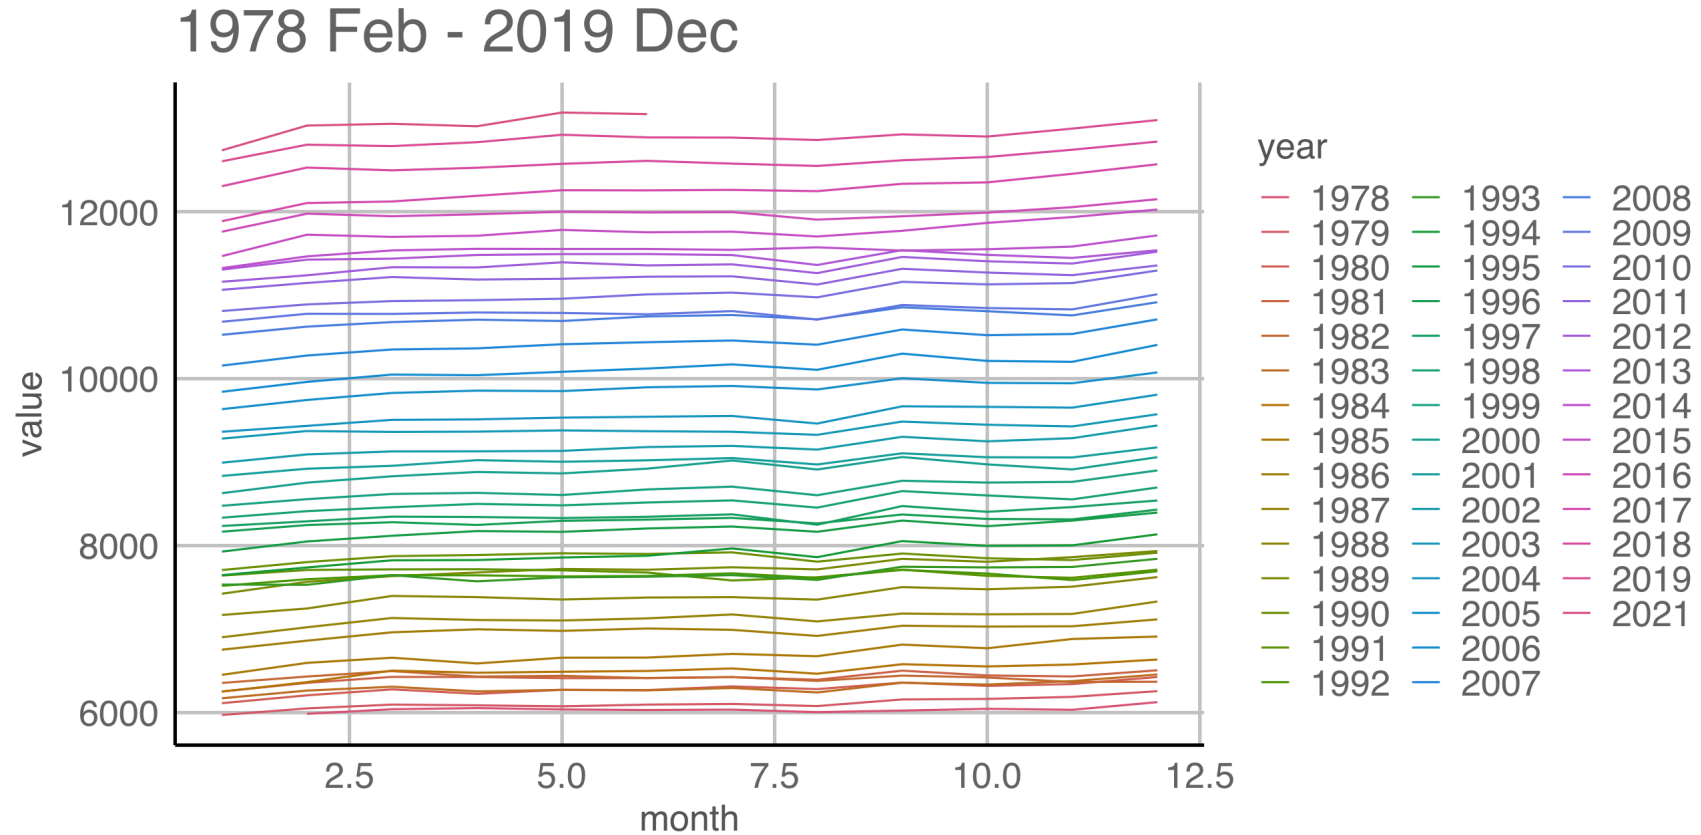
```
## $ month <dbl> 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 1, 2, 3, 4, 5, 6, 7, 8, 9,
```

```
## $ year <fct> 1978, 1978, 1978, 1978, 1978, 1978, 1978, 1978, 1978, 1978, 197
```

```
## $ value <dbl> 5985.660, 6040.561, 6054.214, 6038.265, 6031.342, 6036.084, 600
```


Case study 2 Employment Data in Australia Part 2/3

Do you notice anything?




Why do you think the number of people employed is going up each year?

Case study 2 Employment Data in Australia Part 3/3



- There's a suspicious change in August numbers from 2014.



- A potential explanation for this is that there was a *change in the survey from 2014*.

**Check if the *data collection*
method has been consistent**

Example 3 Experimental layout and data Part 1/2



lecture3-example3.csv

```
df3 <- read_csv(here::here("data/lecture3-example3.csv"),
                col_types = cols(
                  row = col_factor(),
                  col = col_factor(),
                  yield = col_double(),
                  trt = col_factor(),
                  block = col_factor()))
```


```
skimr::skim(df3)
```

```
## — Data Summary —————
##                               Values
## Name                         df3
## Number of rows               48
## Number of columns            5
## -----
## Column type frequency:
```

Example 3 Experimental layout and data Part 2/2



- The experiment tests the effects of 9 fertilizer treatments on the yield of brussel sprouts on a field laid out in a rectangular array of 6 rows and 8 columns.



- High sulphur and high manure seems to best for the yield of brussel sprouts.
- Any issues here?

Check if experimental layout given in the data and the description match

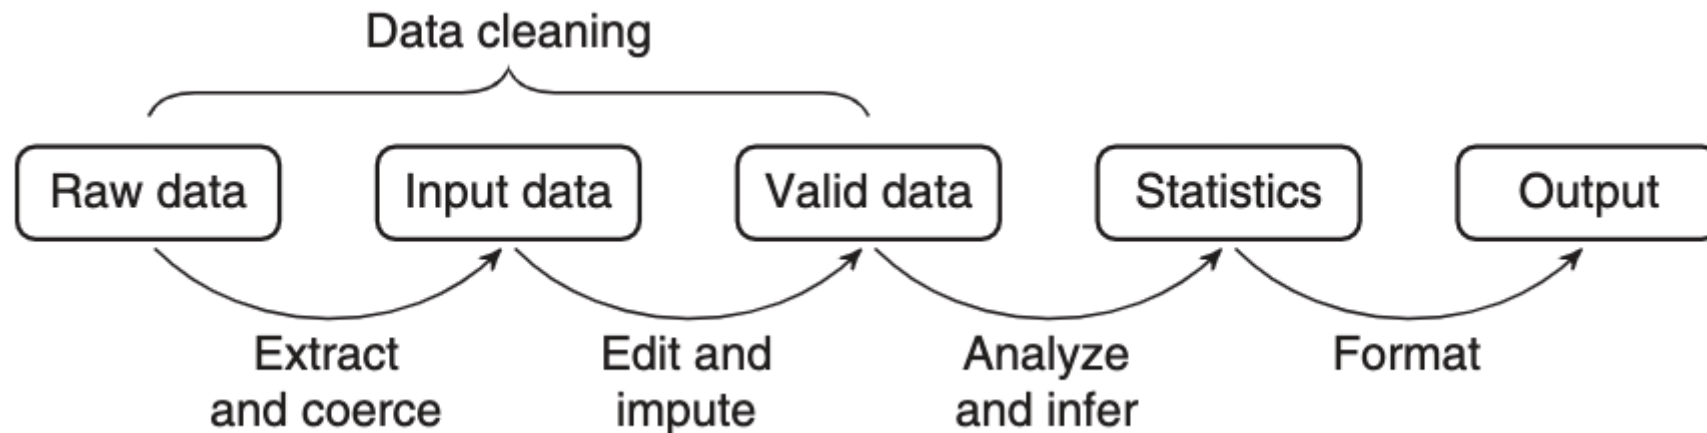
In particular, have a check with a plot to see if treatments are *randomised*.

Statistical Value Chain

“

*... a **statistical value chain** is constructed by defining a number of meaningful intermediate data products, for which a chosen set of quality attributes are well described ...*

— van der Loo & de Jonge (2018)



Case study 3 Dutch supermarket revenue and cost Part 1/3

- Data contains the revenue and cost (in Euros) for 60 supermarkets
- Data has been anonymised and distorted

```
## Rows: 60
## Columns: 11
## $ id          <fct> RET01, RET02, RET03, RET04, RET05, RET06, RET07, RET08, R
## $ size        <fct> sc0, sc3, sc3, sc3, sc3, sc0, sc3, sc1, sc3, sc2, sc2, sc
## $ incl.prob    <dbl> 0.02, 0.14, 0.14, 0.14, 0.14, 0.02, 0.14, 0.02, 0.14, 0.0
## $ staff        <int> 75, 9, NA, NA, NA, 1, 5, 3, 6, 5, 5, 5, 13, NA, 3, 52, 10
## $ turnover     <int> NA, 1607, 6886, 3861, NA, 25, NA, 404, 2596, NA, 645, 287
## $ other.rev     <int> NA, NA, -33, 13, 37, NA, NA, 13, NA, NA, NA, NA, 12, NA,
## $ total.rev     <int> 1130, 1607, 6919, 3874, 5602, 25, 1335, 417, 2596, NA, 64
## $ staff.costs   <int> NA, 131, 324, 290, 314, NA, 135, NA, 147, NA, 130, 182, 3
## $ total.costs   <int> 18915, 1544, 6493, 3600, 5530, 22, 136, 342, 2486, NA, 63
## $ profit        <int> 20045, 63, 426, 274, 72, 3, 1, 75, 110, NA, 9, 220, 34, 8
## $ vat           <int> NA, NA, NA, NA, NA, NA, 1346, NA, NA, NA, NA, NA, NA, 863
```

Case study 3 Dutch supermarket revenue and cost Part 2/3

- Checking for completeness of records

```
library(validate)
rules <- validator(
  is_complete(id),
  is_complete(id, turnover),
  is_complete(id, turnover, profit))
out <- confront(SBS2000, rules)
summary(out)
```

##	name	items	passes	fails	nNA	error	warning	expression
## 1	V1	60	60	0	0	FALSE	FALSE	is_complete(id)
## 2	V2	60	56	4	0	FALSE	FALSE	is_complete(id, turnover)
## 3	V3	60	52	8	0	FALSE	FALSE	is_complete(id, turnover, profit)

Case study 3 Dutch supermarket revenue and cost Part 3/3

- Sanity check derived variables

```
library(validate)
rules <- validator(
  total.rev - profit == total.costs,
  turnover + other.rev == total.rev,
  profit <= 0.6 * total.rev
)
out <- confront(SBS2000, rules)
summary(out)
```

##		name	items	passes	fails	nNA	error	warning	
## 1	V1	60	39	14	7	FALSE	FALSE	abs(total.rev - profit - total.co	
## 2	V2	60	19	4	37	FALSE	FALSE	abs(turnover + other.rev - total.	
## 3	V3	60	49	6	5	FALSE	FALSE	(profit - 0.6 * total.r	

Take away messages

- ✈️ Sanity check your data:
 - by validating the variable types
 - with independent or external sources
 - by checking the data quality
- ✈️ Check if the data collection method has been consistent
- ✈️ Check if experimental layout given in the data and the description match
- ✈️ Consider if or how data were derived for further sanity check of your data

Next we'll have a look at the

2 Model formulation



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).

Lecturer: *Emi Tanaka*

✉ ETC5521.Clayton-x@monash.edu

📅 Week 3 - Session 1